

A Case Study in Functional Conversion and Mode Inference in miniKanren

Ekaterina Verbitskaia
kajigor@gmail.com
JetBrains Research
Belgrade, Serbia

Igor Engel
igorenge@mail.ru
JetBrains Research
Bremen, Germany

Daniil Berezun
daniil.berezun@jetbrains.com
JetBrains Research
Amsterdam, Netherlands

Abstract

One of the most attractive applications of relational programming is inverse computation. It offers an approach to solving complex problems by transforming verifiers into solvers with relatively low effort. Unfortunately, inverse computation often suffers from interpretation overhead, leading to subpar performance compared to direct program inversion. A prior study introduced a functional conversion scheme capable of creating inversions of MINIKANREN specifications with respect to a known fixed direction. This paper expands upon it by providing a semi-automated functional conversion algorithm. Our evaluation demonstrates a significant performance improvement achieved through functional conversion.

CCS Concepts: • Software and its engineering → Constraint and logic languages.

Keywords: program inversion, inverse computations, relational programming, functional programming, conversion

1 Introduction

There is a well-known observation [1, 10] that programs solving certain problems can be acquired by inverting programs solving some other, much simpler, problems. Sometimes the difference in the “simplicity” can be characterized in precise complexity-theoretic terms: for example, type checking for simple typed lambda calculus (STLC) is known to be linear-time (and rather straightforward to implement), while type inference (its inversion) is PTIME-complete [11], and type inhabitation problem (its another inversion) is PSPACE-complete [16].

In the scope of this paper we will be interested in a more concrete scenario of this generic idea, namely, of turning verifiers into solvers. A verifier is a procedure which, given an instance of the problem and some sample, verifies if this sample is a solution. A solver takes an instance of the problem and returns such a sample which makes the verifier to succeed. For the variety of search problems the implementation of a verifier is straightforward; on the other hand its inversion is a solver, which as a rule is much harder to implement in an explicit manner. There are a few approaches to program inversion [2, 3], and the properties of the solver produced by inversion greatly depend on the approach utilized.

We focus on the application of relational programming [5] as a way to run programs in the reverse direction.

Relational programming can be considered as a subfield of conventional logic programming focused on study of implementation techniques and applications of purely relational specifications. In a narrow sense, relational programming amounts to writing programs in MINIKANREN¹ — an embedded DSL initially developed for SCHEME and later ported for dozens of other host languages. Based on the same theory of first-order Horn clauses as, for example, PROLOG, MINIKANREN employs a complete interleaving search [6, 13] and discourages the use of extra-logical features such as knowledge of concrete search order, “cuts”, side-effects, efficient, but non-relational arithmetic, etc. In conventional logic programming the specification provided by the end-user usually encodes a certain concrete way to solve a problem. Contrary to this, MINIKANREN shifts the focus onto the specification of the problem itself with no certain hints on how to solve its various instances. This makes the specifications written in MINIKANREN short, elegant and expressive.

It is possible to directly employ the verifier-to-solver approach [7] with MINIKANREN. It has been successfully applied in a few non-trivial projects [8, 9]. On the other hand, many useful optimization techniques cannot be applied for MINIKANREN programs directly since these programs lack an important part of information — the direction under which relational verifier turns into a solver. By taking this information into account, it is possible to make the approach more universally practically applicable.

In this paper we present the results of our exploration in the area of mode inference and functional conversion for MINIKANREN. Mode analysis and inference is a relevant technique for conventional logic programming [4, 12, 14]. A mode can be considered as an implicit specification of a direction in which a relation is intended to be evaluated. Given a user-defined description of modes for (some) relations, mode analysis propagates the mode information through the rest of the logic program thus defining more concrete evaluation strategy for the rest of its relations.

Various notions and concrete approaches are employed for mode analysis in various settings, and we give a survey in Section 3. In our setting we consider user-defined mode specification for the top-level goal as the prescription

¹Website of the MINIKANREN programming language: <http://minikanren.org>

\mathcal{D}_V^N	: $R_n(x_1, \dots, x_n) = \text{Disj}_V, x_i \in V$	normalized relation definition
Disj_V	: $\bigvee (c_1, \dots, c_n), c_i \in \text{Conj}_V$	normal form
Conj_V	: $\bigwedge (g_1, \dots, g_n), g_i \in \text{Base}_V$	normal conjunction
Base_V	: $V \equiv \mathcal{T}_V$	flat unification
	$R_n(x_1, \dots, x_n), x_i \in V, i \neq j \Rightarrow x_i \neq x_j$	flat call

Figure 1. Abstract syntax of MICROKANREN in the normal form

of the direction in which relational specification has to be evaluated to provide a solver for the problem in question. However, such a prescription cannot be directly employed in MINIKANREN as it contradicts the very nature of relational programming. Instead, we accompany mode inference with functional conversion — a transformation which, given a relational specification, a top-level goal, user-defined modes for this goal and the results of mode inference provides a regular functional program which delivers exactly the same answers as the top-level goal being evaluated in the direction prescribed by the modes. In addition, functional conversion can sometimes eliminate the interpretation overhead introduced by MINIKANREN implementation as a shallow DSL: it is capable of replacing unification with pattern-matching, make use of deterministic order of evaluation if such an order is discovered by mode inference, etc.

The contribution of this paper is as follows:

- We reiterate on mode inference for MINIKANREN, specifying concrete requirements specific for both MINIKANREN and our ultimate goal of putting verifier-to-solver idea to work.
- We describe a concrete approach to mode inference which takes the aforementioned requirements into account. As generally mode inference is known to be undecidable we develop a number of heuristics specific to our case.
- We implement both mode inference and functional conversion for a reference MINIKANREN implementation.
- We evaluate our implementation on several benchmarks, to investigate the advantages, drawbacks, and potential areas for improvement in our approach.

The rest of the paper is structured as follows. ...

2 Normal Form Abstract Syntax of MINIKANREN

In this paper we use core MINIKANREN usually referred to as MICROKANREN. In its syntax, a relation is a goal comprised of disjunctions or conjunctions of other goals. A base goal can be either an explicit unification of two terms or a call to a relation. An example program in MINIKANREN is shown in figure 2. It relates triples of Peano natural numbers x, y, z such that $x + y = z$. We use identifiers which start with the upper-case letters as constructors, and identifiers which start with the lower-case letters as variable names.

```

let rec addo x y z =
  (x ≡ 0 ∧ y ≡ z) ∨
  (fresh x1, z1 in
    x ≡ S x1 ∧
    addo x1 y z1 ∧
    z ≡ S z1)

```

Figure 2. Addition relation in MINIKANREN

A superscript ^o denotes a relation name while the keyword **fresh** introduces fresh variables into the scope. To execute a relation, one should provide a query to run, for example the query **run** q (add^o q q (S (S 0))) finds a number which, doubled, is 2 in Peano representation.

To simplify the functional conversion scheme, we consider MINIKANREN relations to be in the superhomogeneous normal form used in the MERCURY programming language [15]. Converting an arbitrary MINIKANREN relation into the normal form is a simple syntactic transformation, which we omit.

In the normal form, a term is either a variable or a constructor application which is flat and linear. Linearity means that arguments of a constructor are distinct variables. To be flat, a term should not contain any nested constructors. Each constructor has a fixed arity n . Below is the abstract syntax of the term language over the set of variables V .

$$\mathcal{T}_V = V \cup \{C_n(x_1, \dots, x_n) \mid x_i \in V; i \neq j \Rightarrow x_i \neq x_j\}$$

Whenever a term which does not adhere to this form is encountered in a unification or as an argument of a call, it is transformed into a conjunction of several unifications, as illustrated by the following examples:

$$\begin{aligned}
 C(x_1, x_2) &\equiv C(C(y_1, y_2), y_3) \Rightarrow \\
 &\Rightarrow x_1 \equiv C(y_1, y_2) \wedge x_2 \equiv y_3 \\
 C(C(x_1, x_2), x_3) &\equiv C(C(y_1, y_2), y_3) \Rightarrow \\
 &\Rightarrow x_1 \equiv y_1 \wedge x_2 \equiv y_2 \wedge x_3 \equiv y_3 \\
 x \equiv C(y, y) &\Rightarrow x \equiv C(y_1, y_2) \wedge y_1 \equiv y_2 \\
 add^o(x, x, z) &\Rightarrow add^o(x_1, x_2, z) \wedge x_1 \equiv x_2
 \end{aligned}$$

Unification in the normal form is restricted to always unify a variable with a term. We also prohibit using disjunctions inside conjunctions. The normalization procedure declares a new relation whenever this is encountered. This is done to limit the number of possible permutations one has to consider when doing the mode inference.

The complete abstract syntax of the MINIKANREN language used in this paper is presented in figure 1.

3 Related Work

A mode generalizes the concept of a direction; this terminology is commonly used in the conventional logic programming community. In its most primitive form, a mode specifies which arguments of a relation will be known at runtime (input) and which are expected to be computed (output). Several logic programming languages have mode systems used for optimizations, with MERCURY² standing out among them. MERCURY is a modern functional-logic programming language with a complicated mode system capable not only of describing directions, but also specifying if a relation in the given mode is deterministic, among other things [12].

Functional logic programming languages such as MERCURY translate the logic subset of the language into a functional programming language much like we do. MERCURY utilizes a strong system of modes and types to guide the compilation. The first version of mode analysis was based on the abstract interpretation [14]. It is still the one used in the implementation. They also researched the use of constraint systems for mode inference [12]. While being more precise, this system proved to be too slow to be used in the compiler.

Mercury's mode system is tight into a type system. Many MINIKANREN languages are embedded into host languages which are not typed and thus we cannot rely on type information in our conversion.

It is hard to compare our implementation with MERCURY because of the difference in semantics. MERCURY uses a system of prescriptive modes which means that the semantics is defined by the mode assigned to a program. We use modes in a prescriptive fashion, and the semantics of a MINIKANREN program is the same set of answers regardless of the order of subgoals. This allows us to compare runtime of differently modded relational programs.

Even though using abstract interpretation is a standard for doing many kinds of program analysis [?], we decided to abstain from using it for mode analysis. This is done in order to simplify the implementation of the functional conversion in the spirit of MINIKANREN language family.

3.1 Modes

Given an annotation for a relation, mode inference determines modes of each variable of the relation. For some modes, conjunctions in the body of a relation may need reordering to ensure that consumers of computed values come after the producers of said values so that a variable is never used before it is bound to some value. In this project, we employed the least complicated mode system, in which variables may only have an in or out mode. A mode maps variables of a relation to a pair of the initial and final instantiations. The

²Website of the MERCURY programming language: <https://mercurylang.org/>

mode in stands for $g \rightarrow g$, while out stands for $f \rightarrow g$. The instantiation f represents an unbound, or free, variable, when no information about its possible values is available. When the variable is known to be ground, its instantiation is g .

In this paper, we call a pair of instantiations a mode of a variable. figure 3 shows examples of the normalized MINIKANREN relations with modes inferred for the forward and backward directions. We use superscript annotation for variables to represent their modes visually. Notice the different order of conjuncts in the bodies of the `addo` relation in different modes.

```

let doubleo xg→g rf→g =
  addo xg→g1 xg→g2 rf→g ∧
  xg→g1 ≡ xg→g2

let rec addo xg→g yg→g zf→g =
  (xg→g ≡ 0 ∧ yg→g ≡ zf→g) ∨
  (xg→g ≡ S xf→g1 ∧
   addo xg→g1 yg→g zf→g ∧
   zf→g ≡ S zg→g1)

```

(a) Forward direction

```

let doubleo xf→g rg→g =
  addo xf→g1 xf→g2 rg→g ∧
  xg→g1 ≡ xg→g2

let rec addo xf→g yf→g zg→g =
  (xf→g ≡ 0 ∧ yf→g ≡ zg→g) ∨
  (zf→g ≡ S zf→g1 ∧
   addo xf→g1 yf→g zg→g1 ∧
   xf→g ≡ S xg→g1)

```

(b) Backward direction

Figure 3. Normalized doubling and addition relations with mode annotations

4 Functional Conversion for MICROKANREN

In this section, we describe the functional conversion algorithm. The reader is encouraged to first read the paper [18] on the topic, which introduces the conversion scheme on a series of examples.

Functional conversion is done for a relation with a concrete fixed direction. The goal is to create a function which computes the same answers as MICROKANREN would, not necessarily in the same order. Since the search in MICROKANREN is complete, both conjuncts and disjuncts can be reordered

freely: interleaving makes sure that no answers would be lost this way. Moreover, the original order of the subgoals is often suboptimal for any direction but the one which the programmer had in mind when they encoded the relation. When the relational conversion is used to create a relation, the order of the subgoals only really suits the forward direction, in which the relation is often not intended to be run (in this case, it is better to run the original function).

The mode inference results in the relational program with all variables annotated by their modes, and all base subgoals ordered in a way that further conversion makes sense. Conversion then produces functions in the intermediate language. It may then be pretty printed into concrete functional programming languages, in our case HASKELL and OCAML.

4.1 Mode Inference

We employ a simple version of mode analysis to order subgoals properly in the given direction. The mode analysis makes sure that a variable is never used before it is associated with some value. It also ensures that once a variable becomes ground, it never becomes free, thus the value of a variable is never lost. The mode inference pseudocode is presented in listing 1.

```

1 modeInfer ( $R_i(x_1, \dots, x_{k_i}) \equiv body$ ) =
2   ( $R_i(x_1, \dots, x_{k_i}) \equiv (\text{modeInferDisj } body)$ )
3
4 modeInferDisj ( $\vee(c_1, \dots, c_n)$ ) =
5    $\vee(\text{modeInferConj } c_1, \dots, \text{modeInferConj } c_n)$ 
6
7 modeInferConj ( $\wedge(g_1, \dots, g_n)$ ) =
8   let (picked, others) = pickConjunct [ $g_1, \dots, g_n$ ]
9   in let moddedPicked = modeInferBase picked
10  in let moddedConjs = modeInferConj ( $\wedge others$ )
11  in  $\wedge(\text{moddedPicked} : \text{moddedConjs})$ 
12
13 pickConjunct goals =
14   pickGuard goals <|>
15   pickAssignment goals <|>
16   pickMatch goals <|>
17   pickCallWithGroundArguments goals <|>
18   pickUnificationGenerator goals <|>
19   pickCallGenerator goals

```

Listing 1. Mode inference pseudocode

Mode inference starts by initializing modes for all variables in the body of the given relation according to the given direction. All variables that are among arguments are annotated with their *in* or *out* modes, while all other variables get only their initial instantiations specified as *f*.

Then the body of the relation is analyzed (see line 2). Since the body is normalized, it can only be a disjunction. Each

disjunct is analyzed independently (see line 5) because no data flow happens between them.

Analyzing conjunctions involves analyzing subgoals and ordering them. Let us first consider mode analysis of unifications and calls, and then circle back to the way we order them. Whenever a base goal is analyzed, all variables in it have some initial instantiation, and some of them also have some final instantiation. Mode analysis of a base goal boils down to making all final instantiations ground.

When analyzing a unification, several situations may occur. Firstly, every variable in the unification can be ground, as in $x^{g \rightarrow g} \equiv O$ or in $y^{g \rightarrow ?} \equiv z^{g \rightarrow ?}$ (here ? is used to denote that a final instantiation is not yet known). We call this case *guard*, since it is equivalent to checking that two values are the same.

The second case is when one side of a unification only contains ground variables. Depending on which side is ground, we call this either *assignment* or *match*. The former corresponds to assigning the value to a variable, as in $x^{f \rightarrow ?} \equiv S x_1^{g \rightarrow g}$ or $x^{g \rightarrow g} \equiv y^{f \rightarrow ?}$. The latter — to pattern matching with the variable as the scrutinee, as in $x^{g \rightarrow g} \equiv S x_1^{f \rightarrow ?}$. Notice that we allow for some variables on the right-hand side to be ground in matches, given that at least one of them is free.

The last case occurs when both the left-hand and right-hand sides contain free variables. This does not translate well into functional code. Any free logic variable corresponds to the possibly infinite number of ground values. To handle this kind of unification, we propose to use *generators* which produce all possible ground values a free variable may have.

We base our ordering strategy for conjuncts on the fact that these four different unification types have different costs. The guards are just equality checks which are inexpensive and can reduce the search space considerably. Assignments and matches are more involved, but they still take much less effort than generators. Moreover, executing non-generator conjuncts first can make some of the variables of the prospective generator ground thus avoiding generation in the end. This is the base reasoning which is behind our ordering strategy.

The function *pickConjunct* selects the base goal which is least likely to blow up the search space. The right-associative function <|> used in lines 14 through 18 is responsible for selecting the base goals in the order described. The function first attempts to pick a base goal with its first argument, and only if it fails, the second argument is called. As a result, *pickConjunct* first picks the first guard unification it can find (*pickGuard*). If no guard is present, then it searches for the first assignment (*pickAssignment*), and then for the match (*pickMatch*). If all unifications in the conjunction are generators, then we search for relation calls with some ground arguments (*pickCallWithGroundArguments*). If there are none, then we have no choice but selecting a

generating unification (`pickUnificationGenerator`) and then a call with all arguments free (`pickCallGenerator`).

Once one conjunct is picked, it is analyzed (see line 9). The picked conjunct may instantiate new variables, thus this information is propagated onto the rest of the conjuncts. Then the rest of the conjuncts is mode analyzed as a new conjunction (see line 10). If any new modes for any of the relations are encountered, they are also mode analyzed.

It is worth noticing that any relation can generate infinitely many answers. We cannot judge the relation to be such generator solely by its mode: for example, the addition relation in the mode $\text{add}^o x^{g \rightarrow g} y^{f \rightarrow g} z^{f \rightarrow g}$ generates an infinite stream, while $\text{add}^o x^{f \rightarrow g} y^{f \rightarrow g} z^{g \rightarrow g}$ does not.

4.2 Conversion into Intermediate Representation

To represent nondeterminism, our functional conversion uses the basis of MICROKANREN — the stream data structure. A relation is converted into a function with n arguments which returns a stream of m -tuples, where n is the number of the input arguments, and m — the number of the output arguments of the relation. Since stream is a monad, functions can be written elegantly in HASKELL using `do`-notation (see figure 5). We use an intermediate representation which draws inspiration from HASKELL's `do`-notation, but can then be pretty-printed into other functional languages. The abstract syntax of our intermediate language is shown in figure 4. The conversion follows quite naturally from the modded relation and the syntax of the intermediate representation.

A body of a function is formed as an interleaving concatenation of streams (**Sum**), each of which is constructed from one of the disjuncts of the relation. A conjunction is translated into a sequence of bind statements (**Bind**): one for each of the conjuncts and a return statement (**Return**) in the end. A bind statement binds a tuple of variables (or nothing) with values taken from the stream in the right-hand side.

A base goal is converted into a guard (**Guard**), match (**Match**), or function call, depending on the goal's type. Assignments are translated into binds with a single return statement on the right. Notice, that a match only has one branch. This branch corresponds to a unification. If the scrutinee does not match the term it is unified with, then an empty stream is returned in the catch-all branch. If a term in the right-hand side of a unification has both out and in variables, then additional guards are placed in the body of the branch to ensure the equality between values bound in the pattern and the actual ground values.

Generators (**Gen**) are used for unifications with free variables on both sides. A generator is a stream of possible values for the free variables, and it is used for each variable from the right-hand side of the unification. The variable from the left-hand side of the unification is then simply assigned the value constructed from the right-hand side. Our current implementation works with an untyped deeply embedded

MICROKANREN, in which there is not enough information to produce generators automatically. We decided to delegate the responsibility to provide generators to the user: a generator for each free variable is added as an argument of the relation. When the user is to call the function, they have to provide the suitable generators.

5 Examples

In this section, we provide some examples which demonstrate mode analysis and conversion results.

5.1 Multiplication Relation

Figure 5 shows the implementation of the multiplication relation mul^o , the mode analysis result for mode $\text{mul}^o x^{f \rightarrow g} y^{g \rightarrow g} z^{g \rightarrow g}$, and the results of functional conversion into HASKELL and OCAML.

Note that the unification comes last in the second disjunct. This is because before the two relation calls are done, both variables in the unification are free. Our version of mode inference puts the relation calls before the unification, but the order of the calls depends on their order in the original relation. There is nothing else our mode inference uses to prefer the order presented in the figure over the opposite: $\text{mul}^o x_1^{f \rightarrow g} y^{g \rightarrow g} z_1^{f \rightarrow g} \wedge \text{add}^o y^{g \rightarrow g} z_1^{g \rightarrow g} z^{g \rightarrow g}$. However, it is possible to derive this optimal order, if determinism analysis is employed: $\text{add}^o y^{g \rightarrow g} z_1^{f \rightarrow g} z^{g \rightarrow g}$ is deterministic while $\text{mul}^o x_1^{f \rightarrow g} y^{g \rightarrow g} z_1^{f \rightarrow g}$ is not. Putting nondeterministic computations first makes the search space larger, and thus should be avoided if another order is possible.

Functional conversions in both languages are similar, modulo the syntax. The HASKELL version employs `do`-notation, while we use `let`-syntax in the OCAML code. Both are syntactic sugar for monadic computations over streams. We use the following convention to name the functions: we add a suffix to the relation's name whose length is the same as the number of the relation's arguments. The suffix consists of the letters I and O which denote whether the argument in the corresponding position is in or out. The function `msum` uses the interleaving function `mplus` to concatenate the list of streams constructed from disjuncts. To check conditions, we use the function `guard` which fails the monadic computation if the condition does not hold. Note that even though patterns for the variable `x0` in the function `addoIOI` are disjunct in two branches, we do not express them as a single pattern match. Doing so would improve readability, but it does not make a difference when it comes to the performance, according to our evaluation.

5.2 The Mode of Addition Relation which Needs a Generator

Consider the example of the addition relation in mode $\text{add}^o x^{g \rightarrow g} y^{f \rightarrow g} z^{f \rightarrow g}$ presented in figure 6. The unification in the first disjunct of this relation involves two free variables.

\mathcal{F}_V	Sum $[\mathcal{F}_V]$	concatenation of streams
	Bind $[[V], \mathcal{F}_V]$	monadic bind for streams
	Return $[\mathcal{T}_V]$	return of a tuple of terms
	Guard (V, V)	equality check
	Match $_V (\mathcal{T}_V, \mathcal{F}_V)$	match a variable against a pattern
	$R_n([V], [G])$	function call
	Gen $_G$	generator

Figure 4. Abstract syntax of the intermediate language \mathcal{F}

```

let rec mulo x y z =
  (x ≡ 0 ∧ z ≡ 0) ∨
  (fresh x1, z1 in
    (x ≡ S x1 ∧
     addo y z1 z ∧
     mulo x1 y z1))

```

(a) Implementation in MINIKANREN

```

muloOII x1 x2 = msum
[ do { let {x0 = 0}
    ; guard (x2 == 0)
    ; return x0 }
, do { x4 ← addIOI x1 x2
    ; x3 ← muloOII x1 x4
    ; let {x0 = S x3}
    ; return x0 } ]
addIOI x0 x2 = msum
[ do { guard (x0 == 0)
    ; let {x1 = x2}
    ; return x1 }
, do { x3 ← case x0 of
    { S y3 → return y3
    ; _ → mzero }
    ; x4 ← case x2 of
    { S y4 → return y4
    ; _ → mzero }
    ; x1 ← addIOI x3 x4
    ; return x1 } ]

```

(c) Functional conversion into HASKELL

```

let rec mulo xf→g yg→g zg→g =
  (xf→g ≡ 0 ∧ zg→g ≡ 0) ∨
  (addo yg→g z1f→g zg→g ∧
   mulo x1f→g yg→g z1g→g ∧
   xf→g ≡ S x1g→g)

```

(b) Mode inference result

```

let rec muloOII x1 x2 = msum
[ ( let* x0 = return 0 in
    let* _ = guard (x2 = 0) in
    return x0 )
; ( let* x4 = addIOI x1 x2 in
    let* x3 = muloOII x1 x4 in
    let* x0 = return (S x3) in
    return x0 ) ]
and addIOI x0 x2 = msum
[ ( let* _ = guard (x0 = 0) in
    let* x1 = return x2 in
    return x1 )
; ( let* x3 = match x0 with
    | S y3 → return y3
    | _ → mzero in
    let* x4 = match x2 with
    | S y4 → return y4
    | _ → mzero in
    let* x1 = addIOI x3 x4 in
    return x1 ) ]

```

(d) Functional conversion into OCAML**Figure 5.** Multiplication relation

We use a generator `gen_addoII0_x2` to generate a stream of ground values for the variable `z` which is passed into the function `addII0` as an argument. It is up to the user to provide a suitable generator. One of the possible generators which produces all Peano numbers in order and an example of its usage are presented in figure 6b.

The generators which produce an infinite stream should be inverse eta-delayed in OCAML and other non-lazy languages. Otherwise, the function would not terminate trying

to eagerly produce all possible ground values before using any of them.

It is possible to automatically produce generators from the data type of a variable, but it is currently not implemented, as we work with an untyped version of MICROKANREN.

6 Evaluation

To evaluate our functional conversion scheme, we implemented the proposed algorithm in HASKELL. We compared

```

let rec addo xg→g yf→g zf→g =
  (xg→g ≡ 0 ∧ yf→g ≡ zf→g) ∨
  (xg→g ≡ S x1f→g ∧
   addo x1g→g yf→g z1f→g ∧
   zf→g ≡ S z1g→g)

```

(a) Mode inference result

```

genNat = msum
[ return 0
, do { x ← genNat
      ; return (S x) } ]
runAddoIO x = addoIO x genNat

```

(b) Generator of Peano numbers

```

addoIO0 x0 gen_addoIO0_x2 = msum
[ do { guard (x0 == 0)
      ; (x1, x2) ← do { x2 ← gen_addoIO0_x2 ; return (x2, x2) }
      ; return (x1, x2) }
, do { x3 ← case x0 of { S y3 → return y3 ; _ → mzero }
      ; (x1, x4) ← addoIO0 x3 gen_addoIO0_x2
      ; let {x2 = S x4} ; return (x1, x2) } ]

```

(c) Functional conversion

Figure 6. Addition relation when only the first argument is in

execution time of several OCANREN relations in different directions against their functional counterparts in the OCAML language. Here we showcase two relational programs and their conversions. The implementation of the functional conversion³ as well as the execution code⁴ can be found on Github.

6.1 Evaluator of Propositional Formulas

In this example, we converted a relational evaluator of propositional formulas: see figure 7. It evaluates a propositional formula *fm* in the environment *st* to get the result *u*. A formula is either a boolean literal, a numbered variable, a negation of another formula, a conjunction or a disjunction of two formulas. Converting it in the direction when everything but the formula is in (see figure 7a), allows one to synthesize formulas which can be evaluated to the given value. The conversion of this relation does not involve any generators and is presented in figure 7b.

We ran an experiment to compare the execution time of the relational interpreter vs. its functional conversion. In the experiment, we generated from 1000 to 10000 formulas which evaluate to true and contain up to 3 variables with known values. The results are presented in figure 8. The functional conversion improved execution time of the query about 2.5 times from 724ms to 291ms for retrieving 10000 formulas.

6.2 Multiplication

In this example, we converted the multiplication relation in several directions and compared them to the relational counterparts: see figure 9. Functional conversion significantly reduced execution time in most directions.

In the forward direction, we run the query $\text{mul}^o \ n \ 10 \ q$ with *n* in the range from 100 to 1000, and the functional conversion was 2 orders of magnitude faster: 927ms vs 9.4ms for the largest *n*, see figure 9a. In the direction which serves as division, we run the query $\text{mul}^o \ (n/10) \ q \ n$ with *n* ranging from 100 to 1000. Here, performance improved 3 orders of magnitude: from 24s to 0.17s for the largest *n*, see figure 9b. Even more impressive was the backward direction $\text{mul}^o \ x^{f→g} \ y^{f→g} \ z^{g→g}$. Querying for all 16 pairs of divisors of 1000 ($\text{mul}^o \ q \ r \ 1000$) took OCANREN about 32.9s, while the functional conversion succeeded in 1.1s.

What was surprising was the mode $\text{mul}^o \ x^{g→g} \ y^{f→g} \ z^{f→g}$. In this case, the functional conversion was not only worse than its relational counterpart, its performance degraded exponentially with the number of answers asked. It took almost 1450ms to find the first 7 pairs of numbers *q* and *r* such that $10 * q = r$, while OCANREN was able to execute the query in 0.74ms (see figure 9c). The source of this terrible behavior was the suboptimal order of the calls in the second disjunct of the mul^o relation in the corresponding mode (see figure 9d). In this case, the call $\text{add}^o \ y^{f→g} \ z_1^{f→g} \ z^{f→g}$ is put first, which generates all possible triples in the addition relation before filtering them by the call $\text{mul}^o \ x_1^{g→g} \ y^{g→g} \ z_1^{g→g}$. The other order of calls is much better (see figure 9e): it is an order of magnitude faster than its relational source. To achieve the better of these two orders automatically, we delay picking any call with all arguments free. A call of this kind

³The repository of the functional conversion project https://github.com/kajigor/uKanren_transformations

⁴Evaluation code <https://github.com/kajigor/miniKanren-func>

```

661 let rec evalo stg→g fmf→g ug→g =
662   ( fmf→g ≡ Lit ug→g ) ∨
663   ( elemo zf→g stg→g ug→g ∧
664     fmf→g ≡ Var zg→g ) ∨
665   ( noto vf→g ug→g ∧
666     evalo stg→g xf→g vg→g ∧
667     fmf→g ≡ Neg xg→g ) ∨
668   ( oro vf→g wf→g ug→g ∧
669     evalo stg→g xf→g vg→g ∧
670     evalo stg→g yf→g wg→g ∧
671     fmf→g ≡ Disj xg→g yg→g ) ∨
672   ( ando vf→g wf→g ug→g ∧
673     evalo stg→g xf→g vg→g ∧
674     evalo stg→g yf→g wg→g ∧
675     fmf→g ≡ Conj xg→g yg→g ) ∨
676

```

(a) Mode inference result

```

679 evaloIOI x0 x2 = msum
680   [ do { let {x1 = Lit x2}
681     ; return x1 }
682   , do { x7 ← elemoOII x0 x2
683     ; let {x1 = Var x7}
684     ; return x1 }
685   , do { x5 ← notoOI x2
686     ; x3 ← evaloIOI x0 x5
687     ; let {x1 = Neg x3}
688     ; return x1 }
689   , do { (x5, x6) ← oroOII x2
690     ; x3 ← evaloIOI x0 x5
691     ; x4 ← evaloIOI x0 x6
692     ; let {x1 = Disj x3 x4}
693     ; return x1 }
694   , do { (x5, x6) ← andoOII x2
695     ; x3 ← evaloIOI x0 x5
696     ; x4 ← evaloIOI x0 x6
697     ; let {x1 = Conj x3 x4}
698     ; return x1 } ]
699

```

(b) Functional conversion

Figure 7. Evaluator of propositional formulas

always works as a generator of every tuple of values which are in relation. It is a reasonable heuristics to postpone their execution until its arguments become more instantiated.

6.3 Relational Sorting

This program is written in truly relational style. By definition, a sorted list has its smallest element in its head followed by a sorted list. The implementation of the `sorto` corresponds to this definition literally: see figure 10.




Figure 8. Execution time of the evaluators of propositional formulas, eval [true ; false ; true] q true


This relation can be used for both sorting a list and generating permutations, depending on which argument is passed into it. One drawback this implementation has is that its performance in the two directions is drastically different and hinges on the order of two relation calls to `smallesto` and `sorto`. When the call to `smallesto` comes first, sorting works fine while permutation generation times out on lists of length 4. Reordering two calls makes it possible to generate permutations for longer lists, however sorting direction starts to time out on lists of length 5.

The only way a programmer can implement the relation in such a way that both directions work well, is by duplicating a conjunction with the two orders mentioned. Even though it leads to somewhat decent performance, it is far from elegant and also increases the amount of work to be done to compute any answer. Mode analysis is a better approach to reordering the conjuncts according to the direction needed. Accompanied by the functional conversion, it also improves the performance significantly: see table 11.


It is worth noting that this relation executes too slowly to be practical even after the functional conversion. It comes from the properties of the algorithm as well as using Peano numbers. However this relation is illustrative of the ways relational programs are supposed to be written and that their execution in the reverse direction can be improved by using sophisticated analyses rather than resorting to inelegant software engineering practices.

6.4 Deterministic Directions


Running in some directions, relations produce deterministic results. For example, any forward direction of a relation



(a) Multiplication: mulo n 10 q



(b) Division: mulo (n/10) q n



(c) Generation: take n (mulo 10 q r)

```

let rec mulo xg→g yf→g zf→g =
  (xg→g ≡ 0 ∧ zf→g ≡ 0) ∨
  (xg→g ≡ S x1f→g ∧
   addo yf→g z1f→g zf→g ∧
   mulo x1g→g yg→g z1g→g )

```

(d) Inefficient mode

```

let rec mulo xg→g yf→g zf→g =
  (xg→g ≡ 0 ∧ zf→g ≡ 0) ∨
  (xg→g ≡ S x1f→g) ∧
  mulo x1g→g yf→g z1f→g ∧
  addo yg→g z1g→g zf→g )

```

(e) Efficient mode

Figure 9. Execution times of the multiplication relation

created by the relational conversion produces a single result, since it mimics the original function. The guard directions are semi-deterministic: they may fail, but if they succeed, they produce a single unit value. If the addition relation is run with one of the first two arguments out, it acts as subtraction and is also deterministic.

For such directions, there is no need to model nondeterminism with the Stream monad. Semi-determinism can be expressed with a Maybe monad, while deterministic directions

can be converted into simple functions. Our implementation of functional conversion only restricts the computations to be monadic, it does not specify which monad to use. By picking other monads, we can achieve performance improvement. For example, using Maybe for division reduces its execution time 30 times in addition to the 2 orders of magnitude improvement from the functional conversion itself: see figure ??

```

let rec sorto xs sorted =
  (xs ≡ [] ∧ sorted ≡ []) ∨
  (fresh smallest others sorted1 in
    xs ≡ smallest : sorted1 ∧
    sorto others sorted1 ∧
    smallesto xs smallest others)

```

Figure 10. Relational sorting in MINIKANREN

7 Discussion

Our experiments indicated that the functional conversion is capable of improving performance of relational computations significantly in the known directions. The improvement stems from eliminating costly unifications in favor of the cheaper equality checks and pattern matches. Besides this, we employed some heuristics which push lower-cost computations to happen sooner while delaying higher-cost ones. It is also possible to take into account determinism of some directions and improve performance of them even more by picking an appropriate monad.

It is not currently clear if the heuristics we used are universal enough. However, it is always safe to run any deterministic computations because they never increase the search space. We believe that it is necessary to integrate determinism check in the mode analysis so that the more efficient modes such as the one presented in figure 9e could be achieved more justifiably.

We also think that further integration with specialization techniques such as partial deduction may benefit the conversion even more [17]. For example, the third argument of the propositional evaluator can be either **true** or **false**. Specializing the evaluator for these two values may help to shave off even more time.

8 Conclusion and Future Work

In this paper, we described a semi-automatic functional conversion of a MICROKANREN relation with a fixed direction into a functional language. We implemented the proposed conversion and applied it to a set of relations, resulting in significant performance enhancement, as demonstrated in our evaluation. As part of the future work, we plan to augment the mode analysis with a determinism check. We also plan to integrate the functional conversion with specialization techniques such as partial deduction.

References

- [1] Sergei Abramov and Robert Glück. 2000. Combining Semantics with Non-standard Interpreter Hierarchies. In *FST TCS 2000: Foundations of Software Technology and Theoretical Computer Science*, Sanjiv Kapoor and Sanjiva Prasad (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 201–213.
- [2] Sergei Abramov and Robert Glück. 2002. *Principles of Inverse Computation and the Universal Resolving Algorithm*. Springer Berlin

- Heidelberg, Berlin, Heidelberg, 269–295. https://doi.org/10.1007/3-540-36377-7_13
- [3] Bogdan Aman, Gabriel Ciobanu, Robert Glück, Robin Kaarsgaard, Jarkko Kari, Martin Kutrib, Ivan Lanese, Claudio Antares Mezzina, Łukasz Mikulski, Rajagopal Nagarajan, et al. 2020. Foundations of reversible computation. *Reversible Computation: Extending Horizons of Computing: Selected Results of the COST Action IC1405 12* (2020), 1–40.
- [4] Saumya K Debray and David S Warren. 1988. Automatic mode inference for logic programs. *The Journal of Logic Programming* 5, 3 (1988), 207–229.
- [5] Daniel P. Friedman, William E. Byrd, and Oleg Kiselyov. 2005. *The Reasoned Schemer*. The MIT Press. <https://doi.org/10.7551/mitpress/5801.001.0001>
- [6] Oleg Kiselyov, Chung-chieh Shan, Daniel P. Friedman, and Amr Sabry. 2005. Backtracking, Interleaving, and Terminating Monad Transformers: (Functional Pearl). In *Proceedings of the Tenth ACM SIGPLAN International Conference on Functional Programming* (Tallinn, Estonia) (ICFP '05). Association for Computing Machinery, New York, NY, USA, 192–203. <https://doi.org/10.1145/1086365.1086390>
- [7] Dmitry Kosarev, Petr Lozov, and Dmitry Boulytchev. 2020. Relational synthesis for pattern matching. In *Asian Symposium on Programming Languages and Systems*, Bruno C. d. S. Oliveira (Ed.). Springer, Springer International Publishing, Cham, 293–310.
- [8] Dmitrii Kosarev, Peter Lozov, Denis Fokin, and Dmitry Boulytchev. 2022. On a Declarative Guideline-Directed UI Layout Synthesis. (2022).
- [9] Peter Lozov, Dmitry Kosarev, Dmitry Ivanov, and Dmitry Boulytchev. 2023. Relational Solver for Java Generics Type System. In *Logic-Based Program Synthesis and Transformation*. Springer International Publishing.
- [10] Petr Lozov, Ekaterina Verbitskaia, and Dmitry Boulytchev. 2019. Relational interpreters for search problems. In *Relational Programming Workshop*. 43.
- [11] HARRY G. MAIRSON. 2004. FUNCTIONAL PEARL Linear lambda calculus and PTIME-completeness. *Journal of Functional Programming* 14, 6 (2004), 623–633. <https://doi.org/10.1017/S0956796804005131>
- [12] David Overton, Zoltan Somogyi, and Peter J Stuckey. 2002. Constraint-based mode analysis of Mercury. In *Proceedings of the 4th ACM SIGPLAN international conference on Principles and practice of declarative programming*. 109–120.
- [13] Dmitry Rozplokhas, Andrey Vyatkin, and Dmitry Boulytchev. 2020. Certified Semantics for Relational Programming. In *Asian Symposium on Programming Languages and Systems*, Bruno C. d. S. Oliveira (Ed.). Springer, Springer International Publishing, Cham, 167–185. https://doi.org/10.1007/978-3-030-64437-6_9
- [14] Zoltan Somogyi. 1987. A System of Precise Models for Logic Programs.. In *ICLP*. Citeseer, 769–787.
- [15] Zoltan Somogyi, Fergus Henderson, and Thomas Conway. 1996. The execution algorithm of Mercury, an efficient purely declarative logic programming language. *The Journal of Logic Programming* 29, 1-3 (1996), 17–64.
- [16] Pawel Urzyczyn. 1997. Inhabitation in typed lambda-calculi (a syntactic approach). In *Typed Lambda Calculi and Applications*, Philippe de Groote and J. Roger Hindley (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 373–389.
- [17] Ekaterina Verbitskaia, Daniil Berezun, and Dmitry Boulytchev. 2021. An Empirical Study of Partial Deduction for miniKanren. *arXiv preprint arXiv:2109.02814* (2021).
- [18] Ekaterina Verbitskaia, Daniil Berezun, and Dmitry Boulytchev. 2022. On a Direction-Driven Functional Conversion. In *Relational Programming Workshop*.

	Relation		Function		Relation		Function
	sorto	smallesto			smallesto	sorto	
[3;2;1;0]	0.077s	0.004s	0.000s	[0;1;2]	0.013s	0.004s	0.004s
[4;3;2;1;0]	timeout	0.005s	0.000s	[0;1;2;3]	timeout	0.005s	0.005s
[31;...;0]	timeout	1.058s	0.006s	[0;...;6]	timeout	0.999s	0.021s
[262;...;0]	timeout	timeout	1.045s	[0;...;8]	timeout	timeout	1.543s
(a) Sorting direction				(b) Permutation generation direction			

Figure 11. Relational sorting evaluation results

1101	1156
1102	1157
1103	1158
1104	1159
1105	1160
1106	1161
1107	1162
1108	1163
1109	1164
1110	1165
1111	1166
1112	1167
1113	1168
1114	1169
1115	1170
1116	1171
1117	1172
1118	1173
1119	1174
1120	1175
1121	1176
1122	1177
1123	1178
1124	1179
1125	1180
1126	1181
1127	1182
1128	1183
1129	1184
1130	1185
1131	1186
1132	1187
1133	1188
1134	1189
1135	1190
1136	1191
1137	1192
1138	1193
1139	1194
1140	1195
1141	1196
1142	1197
1143	1198
1144	1199
1145	1200
1146	1201
1147	1202
1148	1203
1149	1204
1150	1205
1151	1206
1152	1207
1153	1208
1154	1209
1155	1210